

```
In [1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sn
from sklearn.cluster import KMeans
import scipy.cluster.hierarchy as sch
from sklearn.cluster import AgglomerativeClustering
```

```
In [2]: data=pd.read_excel('EastWestAirlines.xlsx',sheet_name='data')
data.head()
```

```
Out[2]:
```

	ID#	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	F
0	1	28143	0	1	1	1	174	1		0
1	2	19244	0	1	1	1	215	2		0
2	3	41354	0	1	1	1	4123	4		0
3	4	14776	0	1	1	1	500	1		0
4	5	97752	0	4	1	1	43300	26		2077

```
In [3]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3999 entries, 0 to 3998
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID#                    3999 non-null   int64
1   Balance                3999 non-null   int64
2   Qual_miles             3999 non-null   int64
3   cc1_miles               3999 non-null   int64
4   cc2_miles               3999 non-null   int64
5   cc3_miles               3999 non-null   int64
6   Bonus_miles             3999 non-null   int64
7   Bonus_trans             3999 non-null   int64
8   Flight_miles_12mo       3999 non-null   int64
9   Flight_trans_12         3999 non-null   int64
10  Days_since_enroll       3999 non-null   int64
11  Award?                  3999 non-null   int64
dtypes: int64(12)
memory usage: 375.0 KB
```

```
In [4]: data.shape
```

```
Out[4]: (3999, 12)
```

```
In [5]: air=data.drop(['ID#','Award?'], axis=1)
air
```

Out[5]:

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
0	28143	0	1	1	1	174	1	0	
1	19244	0	1	1	1	215	2	0	
2	41354	0	1	1	1	4123	4	0	
3	14776	0	1	1	1	500	1	0	
4	97752	0	4	1	1	43300	26	2077	
...
3994	18476	0	1	1	1	8525	4	200	
3995	64385	0	1	1	1	981	5	0	
3996	73597	0	3	1	1	25447	8	0	
3997	54899	0	1	1	1	500	1	500	
3998	3016	0	1	1	1	0	0	0	

3999 rows × 10 columns

In [6]:

```
def norm_func(i):
    x = (i-i.min())/(i.max()-i.min())
    return (x)
```

In [7]:

```
df_norm = norm_func(air.iloc[:, :])
df_norm
```

Out[7]:


	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
0	0.016508	0.0	0.00	0.0	0.0	0.000660	0.011628	0.000000	
1	0.011288	0.0	0.00	0.0	0.0	0.000815	0.023256	0.000000	
2	0.024257	0.0	0.00	0.0	0.0	0.015636	0.046512	0.000000	
3	0.008667	0.0	0.00	0.0	0.0	0.001896	0.011628	0.000000	
4	0.057338	0.0	0.75	0.0	0.0	0.164211	0.302326	0.067398	
...
3994	0.010837	0.0	0.00	0.0	0.0	0.032330	0.046512	0.006490	
3995	0.037766	0.0	0.00	0.0	0.0	0.003720	0.058140	0.000000	
3996	0.043169	0.0	0.50	0.0	0.0	0.096505	0.093023	0.000000	
3997	0.032202	0.0	0.00	0.0	0.0	0.001896	0.011628	0.016225	
3998	0.001769	0.0	0.00	0.0	0.0	0.000000	0.000000	0.000000	

3999 rows × 10 columns

In [8]:

```
fig = plt.figure(figsize=(10, 8))
WCSS = []
for i in range(1, 11):
    clf = KMeans(n_clusters=i)
    clf.fit(df_norm)
    WCSS.append(clf.inertia_)
plt.plot(range(1, 11), WCSS)
plt.title('The Elbow Method')
plt.ylabel('WCSS')
plt.xlabel('Number of Clusters')
```

The Elbow Method



In [9]: WCSS

```
Out[9]: [898.2705822007467,  
         448.9989194445697,  
         331.4315185424921,  
         272.39254909046446,  
         241.07572112696283,  
         217.7640341949783,  
         200.21286104789934,  
         182.7151065322523,  
         172.10646043690986,  
         160.7291592574553]
```

```
In [10]: clf = KMeans(n_clusters=4)  
         y_kmeans = clf.fit_predict(df_norm)
```

In [11]: y_kmeans

```
Out[11]: array([2, 2, 2, ..., 1, 0, 0])
```

In [12]: `clf.cluster_centers_`

```
Out[12]: array([[2.36724363e-02, 1.09788344e-02, 2.37926136e-02, 1.10085227e-02,
1.24289773e-03, 1.69214392e-02, 8.35177722e-02, 1.28437586e-02,
2.19500429e-02, 2.65314415e-01],
[5.01893530e-02, 1.09101606e-02, 6.26102941e-01, 7.35294118e-04,
3.30882353e-03, 1.18152084e-01, 2.09917921e-01, 1.28678021e-02,
2.18091010e-02, 3.63864204e-01],
[3.87082781e-02, 1.52707004e-02, 3.21888412e-02, 1.07296137e-02,
1.50214592e-03, 1.96348371e-02, 9.29234455e-02, 1.58009047e-02,
2.81156369e-02, 6.93991623e-01],
[8.05499702e-02, 1.47843877e-02, 7.54021448e-01, 6.70241287e-04,
8.71313673e-03, 1.78247067e-01, 2.29082237e-01, 1.93799495e-02,
3.37144013e-02, 7.44413255e-01]])
```

```
In [13]: clf.inertia_
```

```
Out[13]: 272.3927662116429
```

```
In [14]: md=pd.Series(y_kmeans)
air['clust']=md
```

```
Out[14]:
```

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Flight_miles_12mo_max
0	28143	0	1	1	1	174	1	0	0
1	19244	0	1	1	1	215	2	0	0
2	41354	0	1	1	1	4123	4	0	0
3	14776	0	1	1	1	500	1	0	0
4	97752	0	4	1	1	43300	26	2077	2077
...
3994	18476	0	1	1	1	8525	4	200	200
3995	64385	0	1	1	1	981	5	0	0
3996	73597	0	3	1	1	25447	8	0	0
3997	54899	0	1	1	1	500	1	500	500
3998	3016	0	1	1	1	0	0	0	0

3999 rows × 11 columns

```
In [15]: air.groupby(air.clust).mean()
```


```
Out[15]:
```

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Flight_miles_12mo_max
clust									
0	40357.669034	122.392045	1.095170	1.022017	1.004972	4461.929688	7.182528	395.8061	395.8061
1	85554.287187	121.805596	3.503682	1.001473	1.013255	31160.790869	18.055965	397.1310	397.1310
2	65991.343348	170.237768	1.128755	1.021459	1.006009	5177.412017	7.991416	486.9364	486.9364
3	137264.839357	164.595716	4.016064	1.001339	1.034806	46974.539491	19.696118	596.4323	596.4323

```
In [16]: plt.figure(figsize=(15,8))
sn.scatterplot(air['clust'],data['ID#'],c=clf.labels_,s=300,marker='*')
plt.show();
```

C:\Users\ROHIT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



```
In [17]: from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
```

```
In [18]: air
```

```
Out[18]:
```

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
0	28143	0	1	1	1	174	1	0	
1	19244	0	1	1	1	215	2	0	
2	41354	0	1	1	1	4123	4	0	
3	14776	0	1	1	1	500	1	0	
4	97752	0	4	1	1	43300	26	2077	
...	
3994	18476	0	1	1	1	8525	4	200	
3995	64385	0	1	1	1	981	5	0	
3996	73597	0	3	1	1	25447	8	0	
3997	54899	0	1	1	1	500	1	500	
3998	3016	0	1	1	1	0	0	0	

3999 rows × 11 columns

```
In [19]: array=air.values
array
```

```
Out[19]: array([[28143,      0,      1, ...,      0, 7000,      2],
               [19244,      0,      1, ...,      0, 6968,      2],
               [41354,      0,      1, ...,      0, 7034,      2],
               ...,
               [73597,      0,      3, ...,      0, 1402,      1],
               [54899,      0,      1, ...,      1, 1401,      0],
               [ 3016,      0,      1, ...,      0, 1398,      0]], dtype=int64)
```

```
In [20]: stscaler = StandardScaler().fit(array)
X = stscaler.transform(array)
X
```

```
Out[20]: array([[ -4.51140783e-01, -1.86298687e-01, -7.69578406e-01, ...,
                 -3.62167870e-01,  1.39545434e+00,  6.04380405e-01],
                 [-5.39456874e-01, -1.86298687e-01, -7.69578406e-01, ...,
                 -3.62167870e-01,  1.37995704e+00,  6.04380405e-01],
                 [-3.20031232e-01, -1.86298687e-01, -7.69578406e-01, ...,
                 -3.62167870e-01,  1.41192021e+00,  6.04380405e-01],
                 ...,
                 [-4.29480975e-05, -1.86298687e-01,  6.83121167e-01, ...,
                 -3.62167870e-01, -1.31560393e+00, -2.75138241e-01],
                 [-1.85606976e-01, -1.86298687e-01, -7.69578406e-01, ...,
                 -9.85033311e-02, -1.31608822e+00, -1.15465689e+00],
                 [-7.00507951e-01, -1.86298687e-01, -7.69578406e-01, ...,
                 -3.62167870e-01, -1.31754109e+00, -1.15465689e+00]])
```

```
In [21]: dbscan = DBSCAN(eps=0.70, min_samples=10)
dbscan.fit(X)
```

```
Out[21]: DBSCAN(eps=0.7, min_samples=10)
```

```
In [22]: dbscan.labels_
```

```
Out[22]: array([0, 0, 0, ..., 5, 6, 6], dtype=int64)
```

```
In [23]: c=pd.DataFrame(dbscan.labels_,columns=['cluster'])
```

```
In [24]: c
```

```
Out[24]:
```

	cluster
0	0
1	0
2	0
3	0
4	-1
...	...
3994	6
3995	6
3996	5
3997	6
3998	6

3999 rows × 1 columns

```
In [25]: df = pd.concat([data,c],axis=1)
df
```

Out[25]:

	ID#	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12m
0	1	28143	0	1	1	1	174	1	
1	2	19244	0	1	1	1	215	2	
2	3	41354	0	1	1	1	4123	4	
3	4	14776	0	1	1	1	500	1	
4	5	97752	0	4	1	1	43300	26	20
...	
3994	4017	18476	0	1	1	1	8525	4	20
3995	4018	64385	0	1	1	1	981	5	
3996	4019	73597	0	3	1	1	25447	8	
3997	4020	54899	0	1	1	1	500	1	50
3998	4021	3016	0	1	1	1	0	0	

3999 rows × 13 columns

```
In [26]: d1=dbscan.labels_
d1
```

Out[26]: array([0, 0, 0, ..., 5, 6, 6], dtype=int64)

```
In [27]: import sklearn
sklearn.metrics.silhouette_score(X, d1)
```

Out[27]: 0.0983201597948486

```
In [28]: from sklearn.cluster import KMeans
clf = KMeans(n_clusters=5)
y_kmeans = clf.fit_predict(X)
```

```
In [29]: y_kmeans
```

Out[29]: array([1, 1, 1, ..., 4, 4, 4])

```
In [30]: cl1=pd.DataFrame(y_kmeans,columns=['Kcluster'])
cl1
```

Out[30]:

Kcluster	
0	1
1	1
2	1
3	1
4	3
...	...
3994	4
3995	4
3996	4
3997	4
3998	4

3999 rows × 1 columns

```
In [31]: df1 = pd.concat([df,cl1],axis=1)
df1
```

Out[31]:

	ID#	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12m
0	1	28143	0	1	1	1	174	1	
1	2	19244	0	1	1	1	215	2	
2	3	41354	0	1	1	1	4123	4	
3	4	14776	0	1	1	1	500	1	
4	5	97752	0	4	1	1	43300	26	20
...	
3994	4017	18476	0	1	1	1	8525	4	20
3995	4018	64385	0	1	1	1	981	5	
3996	4019	73597	0	3	1	1	25447	8	
3997	4020	54899	0	1	1	1	500	1	50
3998	4021	3016	0	1	1	1	0	0	

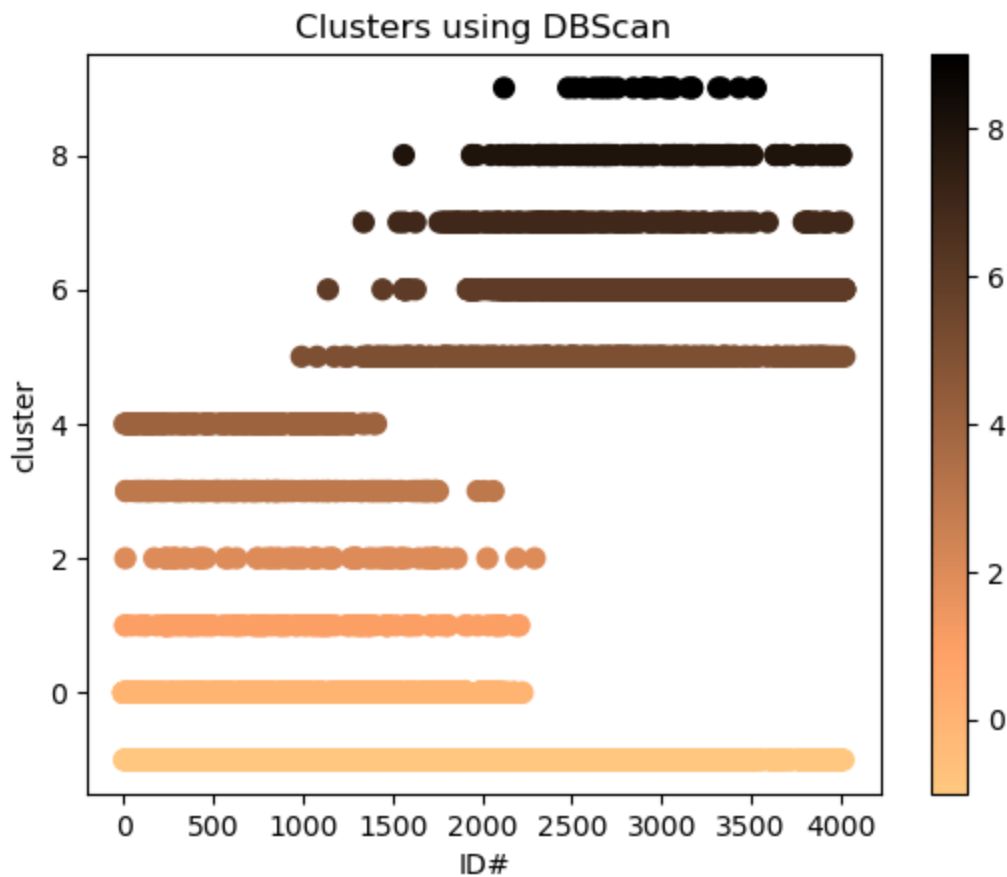
3999 rows × 14 columns

```
In [32]: sklearn.metrics.silhouette_score(X, y_kmeans)
```

Out[32]: 0.30326300970492936

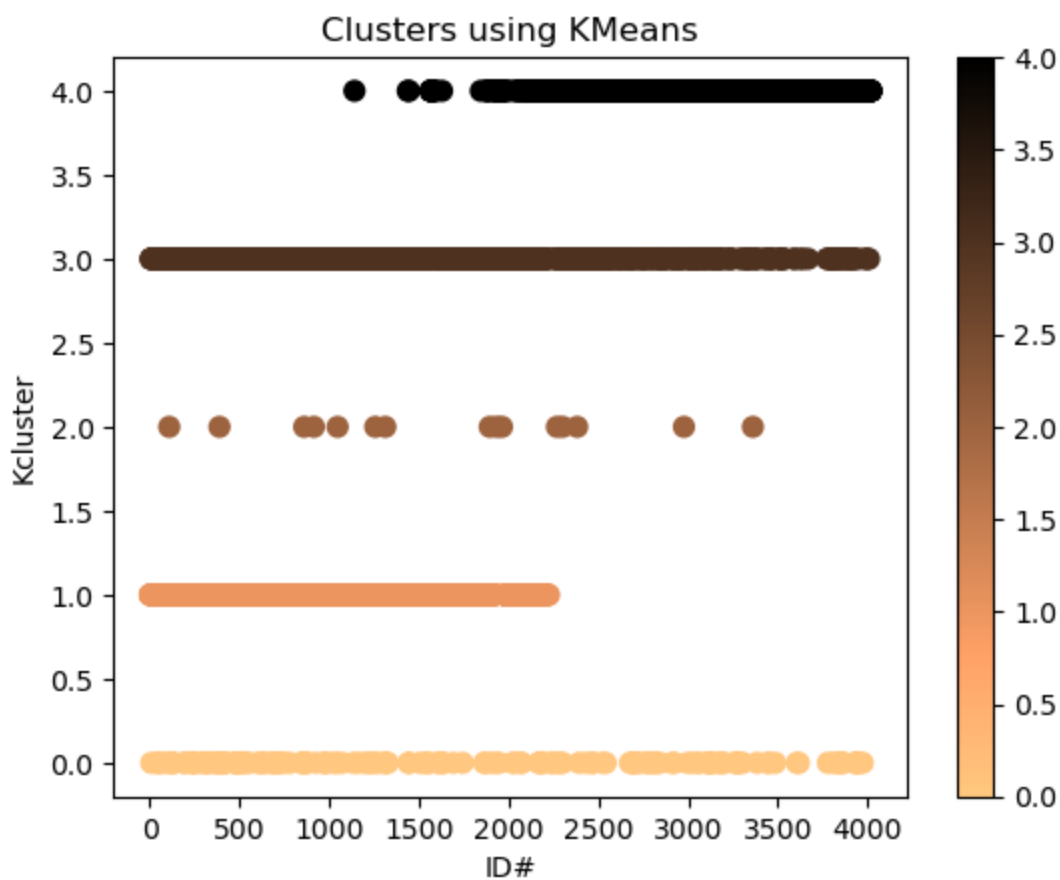
```
In [33]: df.plot(x="ID#",y ="cluster",c=dbscan.labels_ ,kind="scatter",s=50 ,cmap=plt.cm.copper_r
plt.title('Clusters using DBScan')
plt.xlabel("ID#")
plt.ylabel("cluster")
```

Out[33]: Text(0, 0.5, 'cluster')



```
In [34]: df1.plot(x="ID#",y ="Kcluster",c=y_kmeans ,kind="scatter",s=50 ,cmap=plt.cm.copper_r)
plt.title('Clusters using KMeans')
```


```
Out[34]: Text(0.5, 1.0, 'Clusters using KMeans')
```



```
In [35]: plt.figure(figsize=(20,10))
plt.scatter(df1['ID#'],df1['Kcluster'],c=df1['Kcluster'],s=400,marker='*')
```

```
plt.show();
```

```
C:\Users\ROHIT\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pas
s the following variables as keyword args: x, y. From version 0.12, the only valid posit
ional argument will be `data`, and passing other arguments without an explicit keyword w
ill result in an error or misinterpretation.
  warnings.warn(
```



```
In [36]: data
```

Out[36]:	ID#	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12m
0	1	28143	0	1	1	1	174	1	
1	2	19244	0	1	1	1	215	2	
2	3	41354	0	1	1	1	4123	4	
3	4	14776	0	1	1	1	500	1	
4	5	97752	0	4	1	1	43300	26	20
...	
3994	4017	18476	0	1	1	1	8525	4	21
3995	4018	64385	0	1	1	1	981	5	
3996	4019	73597	0	3	1	1	25447	8	
3997	4020	54899	0	1	1	1	500	1	51
3998	4021	3016	0	1	1	1	0	0	

3999 rows × 12 columns

```
In [37]: air=data.drop(['ID#','Award?'],axis=1)
```

```
In [38]: air
```

Out[38]:	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
0	28143	0	1	1	1	174	1	0	
1	19244	0	1	1	1	215	2	0	
2	41354	0	1	1	1	4123	4	0	
3	14776	0	1	1	1	500	1	0	
4	97752	0	4	1	1	43300	26	2077	
...
3994	18476	0	1	1	1	8525	4	200	
3995	64385	0	1	1	1	981	5	0	
3996	73597	0	3	1	1	25447	8	0	
3997	54899	0	1	1	1	500	1	500	
3998	3016	0	1	1	1	0	0	0	


3999 rows × 10 columns

```
In [39]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
air_subset = pd.DataFrame(scaler.fit_transform(air.iloc[:,1:7]))
air_subset
```


Out[39]:	0	1	2	3	4	5
0	-0.186299	-0.769578	-0.098242	-0.062767	-0.702786	-1.104065
1	-0.186299	-0.769578	-0.098242	-0.062767	-0.701088	-0.999926
2	-0.186299	-0.769578	-0.098242	-0.062767	-0.539253	-0.791649
3	-0.186299	-0.769578	-0.098242	-0.062767	-0.689286	-1.104065
4	-0.186299	1.409471	-0.098242	-0.062767	1.083121	1.499394
...
3994	-0.186299	-0.769578	-0.098242	-0.062767	-0.356960	-0.791649
3995	-0.186299	-0.769578	-0.098242	-0.062767	-0.669367	-0.687511
3996	-0.186299	0.683121	-0.098242	-0.062767	0.343804	-0.375096
3997	-0.186299	-0.769578	-0.098242	-0.062767	-0.689286	-1.104065
3998	-0.186299	-0.769578	-0.098242	-0.062767	-0.709992	-1.208203

3999 rows × 6 columns


```
In [50]: from scipy.cluster.hierarchy import linkage
import scipy.cluster.hierarchy as sch # for creating dendrogram
p = np.array(df_norm) # converting into numpy array format
z = linkage(df_norm, method="single", metric="euclidean")
plt.figure(figsize=(15, 5))
plt.title('Hierarchical Clustering Dendrogram')
plt.xlabel('Index')
plt.ylabel('Distance')
sch.dendrogram(z,)
plt.show()
```



```
In [51]: p = np.array(df_norm)
z = linkage(df_norm, method="complete", metric="euclidean")
plt.figure(figsize=(15, 5))
plt.title('Hierarchical Clustering Dendrogram')
plt.xlabel('Index')
plt.ylabel('Distance')
sch.dendrogram(z,)
plt.show()
```



```
In [52]: p = np.array(df_norm)
z = linkage(df_norm, method="average", metric="euclidean")
plt.figure(figsize=(15, 5))
plt.title('Hierarchical Clustering Dendrogram')
plt.xlabel('Index')
plt.ylabel('Distance')
sch.dendrogram(z,)
plt.show()
```



```
In [53]: from sklearn.cluster import AgglomerativeClustering
h_complete = AgglomerativeClustering(n_clusters=5, linkage='complete',affinity = "euclidean")
cluster_labels=pd.Series(h_complete.labels_)
cluster_labels
air['clust']=cluster_labels
air
```

Out[53]:

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Flight_trans_12mo
0	28143	0	1	1	1	174	1	0	0
1	19244	0	1	1	1	215	2	0	0
2	41354	0	1	1	1	4123	4	0	0
3	14776	0	1	1	1	500	1	0	0
4	97752	0	4	1	1	43300	26	2077	0
...
3994	18476	0	1	1	1	8525	4	200	0
3995	64385	0	1	1	1	981	5	0	0
3996	73597	0	3	1	1	25447	8	0	0
3997	54899	0	1	1	1	500	1	500	0
3998	3016	0	1	1	1	0	0	0	0

3999 rows × 11 columns

```
In [54]: air.iloc[:,1:].groupby(air.clust).mean()
```

Out[54]:

	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Flight_trans_12mo
clust								
0	149.668605	3.667151	1.001453	1.001453	38389.199128	18.653343	477.122093	1.424
1	0.000000	3.642857	1.000000	4.142857	97132.785714	28.214286	378.571429	1.070
2	137.437211	1.201079	1.021572	1.001156	5388.251541	7.672188	414.703005	1.240
3	347.000000	2.500000	1.000000	1.000000	65634.250000	69.250000	19960.000000	49.250
4	1355.000000	1.222222	1.000000	1.000000	14267.222222	15.555556	2392.666667	9.111

```
In [55]: data = air[(air.clust==0)]
data
```

Out[55]:

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
4	97752	0	4	1	1	43300	26	2077	
8	443003	0	3	2	1	1753	43	3850	
11	96522	0	5	1	1	61105	19	0	
15	28495	0	4	1	1	49442	15	0	
16	51890	0	4	1	1	48963	16	0	
...
3979	57793	0	3	1	1	20959	15	1198	
3980	28867	0	3	1	1	19169	28	0	
3985	59017	0	4	1	1	34746	25	0	
3991	39142	0	3	1	1	14981	28	0	
3996	73597	0	3	1	1	25447	8	0	

1376 rows × 11 columns

```
In [56]: data = air[(air.clust==1)]
data
```

Out[56]:

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
384	28193	0	5	1	4	103456	32	0	
850	25629	0	4	1	5	107308	36	600	
904	15835	0	2	1	5	112171	16	0	
1036	154664	0	4	1	3	73881	33	0	
1244	166465	0	4	1	4	98717	30	0	
1300	737514	0	3	1	4	43993	40	2150	
1878	128302	0	5	1	4	132263	32	400	
1917	82981	0	5	1	3	105325	44	1000	
1946	126873	0	4	1	5	95598	32	500	
2251	125679	0	3	1	4	68151	27	650	
2283	128456	0	5	1	3	70125	28	0	
2364	108081	0	5	1	5	217006	24	0	
2955	61474	0	1	1	4	47717	14	0	
3338	84148	0	1	1	5	84148	7	0	

```
In [57]: data = air[(air.clust==2)]
data
```

Out[57]:

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
0	28143	0	1	1	1	174	1	0	
1	19244	0	1	1	1	215	2	0	
2	41354	0	1	1	1	4123	4	0	
3	14776	0	1	1	1	500	1	0	
5	16420	0	1	1	1	0	0	0	
...
3993	3974	0	1	1	1	365	3	0	
3994	18476	0	1	1	1	8525	4	200	
3995	64385	0	1	1	1	981	5	0	
3997	54899	0	1	1	1	500	1	500	
3998	3016	0	1	1	1	0	0	0	

2596 rows × 11 columns

In [58]:

data=air[(air.clust==3)]
data

Out[58]:

	Balance	Qual_miles	cc1_miles	cc2_miles	cc3_miles	Bonus_miles	Bonus_trans	Flight_miles_12mo	Fli
2015	53232	888	4	1	1	80696	65	22100	
3235	287033	0	1	1	1	26161	58	12873	
3583	160114	500	1	1	1	71954	86	30817	
3594	27619	0	4	1	1	83726	68	14050	

In [60]:

data = air[(air.clust==4)]

In []: